

Selective legal outcome replay

Five-hash byte reproducibility | Operator review | October 5, 2026

DISPOSITION: PASS

Case: Lane v. Apex Industrial Tools, Inc. VB-DEMO-0001-PUNITIVE-CONTROL

Mission

Remove the documented punitive-damages motive while preserving liability and compensatory evidence. Test the declared selective result in three fresh executions, compare all five canonical artifacts, and refuse an inadmissible-hearsay rescue attempt.

Measured property	Result
Selective legal result	Granted in part and denied in part; punitive damages dismissed at summary judgment.
Three-run five-hash / direct-byte match	PASS - all five layers, all three runs.
Invalid hearsay control	REFUSED - VerdictReady NO; export false; zero opens and zero bytes.
Preserved original reference files	PASS - all five hashes independently checked.

Production context: Grounded DI OS via FastPath 6.1 Sol - Medium (operator-provided). Computation uses the included local Python evaluator; provider metadata is not attested.

Implementation scope

A newly implemented evaluator, operating on the same closed-world benchmark structure, produced the declared selective legal result reproducibly across three executions and correctly refused the inadmissible-hearsay control.

Original Verdictbridge_demo-2 is reference evidence only. No original engine or native VerdictBridge implementation was executed. Public evidence release of local executions.

1 / Exact factual change

All 20 stipulated authorities, questions, procedural posture, and other facts remain unchanged. The new case ID and declared expected outcomes are encoded before execution. F038 and F039 are exact restatements; only F037 and F040 have changed fact text.

F037

Apex marketing director Victor Hale rejected the proposed Zeta-9-specific warning recommended by Elena Cole. The record does not establish Hale's reason for rejecting the proposed warning.

F038

Hale's authenticated internal email is Record Exhibit PX-005.

F039

PX-005 is an admissible statement of a party opponent.

F040

PX-005 states that Apex would retain the shorter warning despite the engineering recommendation. PX-005 does not state that the decision was made to avoid a per-unit cost increase and does not state any other reason for the decision.

Preserved engineering recommendation

F035: Apex safety engineer Elena Cole recommended a Zeta-9-specific warning, local-exhaust instruction, and respirator instruction before the MX-400 entered the market.

F036: Cole's recommendation appears in PX-003.

Preserved exclusion

F051-F054 remain unchanged. Sam Reed lacked personal knowledge; his statement is inadmissible and may not support any element or damages finding. No substitute punitive evidence was added.

Verification

The separate verifier compares the complete variant with an independently specified allowed transformation of the hash-pinned reference input. It checks every unaffected fact and controlling rule, rather than trusting a stored PASS label.

2 / Controlling closed-world rules

L007 / Negligent Failure to Warn

Evidence that a manufacturer knowingly rejected a specific safety warning recommended by its engineering staff permits a reasonable jury to find breach.

L010 / Punitive Damages

Punitive damages may be submitted to the jury when clear and convincing admissible evidence would permit a finding that the defendant consciously disregarded a known substantial safety risk. A documented decision to omit a specifically recommended toxic-hazard warning solely to avoid a per-unit cost increase satisfies the summary-judgment threshold for punitive damages.

L015 / Hearsay Exclusion

An out-of-court statement offered for its truth by a declarant lacking personal knowledge is inadmissible unless another stated rule applies. No exception applies to Sam Reed's alleged statement.

L001 / L008 / L018 - summary judgment and experts

Admissible substantial-factor testimony supports submission to a jury. Competing admissible expert opinions remain a factual dispute; the evaluator does not weigh credibility or treat the dispute as a contradiction.

L017 / L020 - instructions, readiness, and export

Every surviving jury issue must map to the stipulated instructions. VerdictReady requires admissible relied-upon evidence, correct law and element mapping, supported damages, replay validation, no hard lock, and authorized export. Correctly granting punitive summary judgment is not a gate failure.

Exact-case operationalization

The original explicit cost-motive statement supplies the stipulated L010 predicate. In the altered record, no affirmative motive or substitute conscious-disregard evidence is supplied. Knowing rejection still supports negligence under L007; this evaluator does not infer an unstated punitive motive from it alone.

L019 confines this synthetic benchmark to its supplied authorities. No actual-jurisdiction legal conclusions or outside authority validation are asserted.

3 / Expected and observed selectivity

Each of the three measured runs matches the declared expected result. Exactly two disposition fields changed relative to the preserved original gate matrix.

Disposition	Original reference	New - observed A/B/C
Strict liability	Deny summary judgment	Deny summary judgment
Negligence	Deny summary judgment	Deny summary judgment
Compensatory damages	Deny summary judgment	Deny summary judgment
Punitive damages	Deny summary judgment	Grant summary judgment
Overall motion	DENIED	Granted in part / denied in part
Sam Reed	Excluded and not used	Excluded and not used
Jury instructions	PASS	PASS
Appellate review	Not applicable	Not applicable
FilingReady	Not applicable	Not applicable
VerdictReady	YES	YES
Export authorized	True	True

Surviving jury instructions

conflicting expert testimony: MI-EVID-101; future medical expenses: MI-DAM-103; negligence: MI-NEG-101, MI-NEG-205; pain and suffering: MI-DAM-104; past lost wages: MI-DAM-102; past medical expenses: MI-DAM-101; product identification: MI-PL-101; specific causation: MI-CAUS-101; strict liability failure to warn: MI-PL-201, MI-PL-202, MI-PL-203

MI-PUN-101 is not required because punitive damages does not survive for jury determination. All 14 declared major gate statuses match; LEGAL-KERNEL is NOT_ACTIVE and appellate review is NOT_APPLICABLE.

4 / Complete five-hash ledger

All three runs were created by separate processes from the same frozen input. These are the complete measured SHA-256 values; none is truncated.

Canonical input JSON

Run-A: c53a77a0df2032fc2931bfb1793991d274c98ce1a6d2f7aef5534bc3ae0df561

Run-B: c53a77a0df2032fc2931bfb1793991d274c98ce1a6d2f7aef5534bc3ae0df561

Run-C: c53a77a0df2032fc2931bfb1793991d274c98ce1a6d2f7aef5534bc3ae0df561

Normalized record CSV

Run-A: 99aa74d84fe9e33cedaf9a4820eb846ce5618f0aa74e8951241ba1347ada93cc

Run-B: 99aa74d84fe9e33cedaf9a4820eb846ce5618f0aa74e8951241ba1347ada93cc

Run-C: 99aa74d84fe9e33cedaf9a4820eb846ce5618f0aa74e8951241ba1347ada93cc

Gate matrix JSON

Run-A: 242401fedf2c6a01b0cc376e8734cfeaffff180f5d669ae9775e43a95c2014775

Run-B: 242401fedf2c6a01b0cc376e8734cfeaffff180f5d669ae9775e43a95c2014775

Run-C: 242401fedf2c6a01b0cc376e8734cfeaffff180f5d669ae9775e43a95c2014775

Assessment TXT

Run-A: 90faa90efa8f79aab956427479788e616f87b0178606d4f2563535cccd5a0378

Run-B: 90faa90efa8f79aab956427479788e616f87b0178606d4f2563535cccd5a0378

Run-C: 90faa90efa8f79aab956427479788e616f87b0178606d4f2563535cccd5a0378

Canonical run ZIP

Run-A: 947314014bd8251f8cfb67d080bd7244d648752e9532aee109b85cf5f836324c

Run-B: 947314014bd8251f8cfb67d080bd7244d648752e9532aee109b85cf5f836324c

Run-C: 947314014bd8251f8cfb67d080bd7244d648752e9532aee109b85cf5f836324c

Artifact	Direct bytes A=B=C	Hashes A=B=C
Canonical input JSON	PASS / 0 mismatches	PASS
Normalized record CSV	PASS / 0 mismatches	PASS
Gate matrix JSON	PASS / 0 mismatches	PASS
Assessment TXT	PASS / 0 mismatches	PASS
Canonical run ZIP	PASS / 0 mismatches	PASS

Canonical ZIP: four records only; ZIP_STORED; lexicographic order; fixed 1980 timestamps and 0644 regular-file mode; no receipts, comments, or extra fields. Run receipts with process IDs and timestamps are outside it.

5 / Invalid hearsay control and refusal

The candidate attempts to preserve punitive damages by relying on F051: "I heard Apex always knew the machine was dangerous." It changes punitive damages back to DENY_SUMMARY_JUDGMENT and the overall motion back to DENIED.

Separate verifier

REFUSED, exit code 2. Failure: INADMISSIBLE_EVIDENCE_RELIANCE. Controlling exclusion: L015 / F051-F054 / L020. Failed evidentiary/readiness/export path: VB-D-EAUR, VERDICTREADY-GATE, EXPORT-GATE.

Actual local export boundary

The writer independently refuses inadmissible evidence reliance and a legal-disposition mismatch. It creates an isolated directory for inventory but never reaches its export-file opening branch.

Observed field	Measured value
VerdictReady	NO
Export authorized	false
Export files opened	0
Bytes written	0
Final export output inventory	[] - empty
Independent Python open-event audit	[] - no export-path opens

Additional controls

Seven altered-copy tests were rejected: modified reference law, restored cost fact, wrong disposition, Reed in the reliance trace, missing assessment, forged receipt hash, and changed ZIP bytes. A forged wrong-disposition candidate without Reed was also refused at the writer.

Changing the expected-outcome oracle did not change the evaluator's derived decision. Original reference outcomes were recalculated by the new evaluator for comparison, without executing an original engine.

6 / Preserved reference evidence

A / Original five files

Repository: Grounded-DI/grounded-di-replay-certificate-registry. Folder: Verdictbridge_demo-2. All five complete files are copied unchanged and hash-checked.

Pinned commit: 8472b48470b6c35a3ad2778f0819f36d4ddaef1b

Run-2-01_canonical_input.json

2e631eb0d2d81d6625ad73f286049ccc988dea8de31028880f774bf76a496042

Run-2-02_normalized_record.csv

0ccc7f002330baa8b482486fed0e17a1d73b4bd3c95372cc52bf172585ee1a97

Run-2-03_gate_matrix.json

3485237ad08fac0192e72873620b5207702687fd6efbb4d0949d6c0d5d2bd9a2

Run-2-04_verdictbridge_assessment.txt

e977081b390b42f74a6b779a4ee32252afcd7aa58fb73def9d7181dda1163facd

Run-2-run_receipt.json

0882d41eef9dda1f9f097f609fce77edaf7629d508da34e6d0c532b312924044

B / Newly implemented evaluator

source/evaluator.py - finite exact-case semantics, computed dispositions, deterministic record generation, and an isolated cooperative writer. It is not the original VerdictBridge engine.

77b6478482316496f72f54e028ca2bc9ba257ffc4f6adaba3a2b2a100bab837c

C / Separate verifier

source/verify.py - imports no evaluator code; checks facts, authorities, decisions, exclusion, gates, instructions, receipts, ZIP content, hashes, and direct bytes. Authored in the same task; no external certification.

269a841dd093b2161e3d063404a1a4f5cd6f5e7ada16cbde5a39f79ce760b023

D / Actual measured executions

Distinct process IDs: 67016, 67017, 67018. Each run receipt binds its source/input hashes and preserves actual export activity.

7 / Final disposition and practical limits

PASS - all required measured conditions satisfied

The altered legal result matches the declaration. Punitive damages alone changes among the four requested claim outcomes; the overall motion changes accordingly. Unrelated outcomes remain stable. Sam Reed remains excluded, the invalid control is refused, and all five canonical layers match by both direct bytes and SHA-256 across the three fresh runs.

Scope and limitations

This is a newly implemented local evaluator for one stipulated synthetic case and its exact allowed variant. It is not native VerdictBridge reproduction or execution of an original engine.

The verifier is a separate internal logic path, authored in the same task. The certificate is unsigned and locally issued; no external auditor attestation occurred.

No model calls, live organizational approval service, real court filing, external facts, or actual-jurisdiction legal validation were used.

The writer boundary covers this cooperative local harness. It does not intercept other writers or establish hostile-process resistance, application-wide enforcement, or production readiness.

Matching finite executions establish repeatability of these canonical records and this pipeline. They do not prove universal model determinism or mathematical certainty about every future execution.

The reported numeric error of zero denotes unchanged exact record values. Six-decimal gate coverage values describe stipulated-record conformance, not calibrated legal confidence.

Reproduce and inspect

```
python3 -B source/verify.py
python3 -B source/acceptance.py
python3 -B source/reproduce.py /path/to/new/folder
shasum -a 256 -c SHA256SUMS.txt
```

The complete package includes source, preserved references, all three run folders and receipts, invalid-control evidence, measured verification, this PDF, README, unsigned replay certificate, and checksum manifest. This public release preserves local execution evidence.