

VerdictBridge Demo 3

Fresh local replay

PASS | October 7, 2026 | Local review copy

Exactly the same declared output under the same frozen execution conditions, with receipts showing what actually ran, SHA-256 hashes, and direct byte comparisons. Material changes produce a specified changed result or refusal.

Measured outcome

Three fresh evaluator processes (64149, 64150, 64151) produced all five canonical artifacts with matching SHA-256 hashes and direct byte equality. All fifteen artifacts also match their pinned historical counterparts directly. Zero byte differences were observed. This is a replay of the existing closed-world benchmark, not a new benchmark.

Strict liability, negligence and compensatory damages: DENY_SUMMARY_JUDGMENT. Punitive damages: GRANT_SUMMARY_JUDGMENT. Overall motion: GRANTED_IN_PART_DENIED_IN_PART. Sam Reed: EXCLUDED_AND_NOT_USED. Jury instruction coverage: PASS. VerdictReady: YES. Export authorization: true.

Fresh refusal evidence

The candidate relying on Sam Reed to preserve punitive damages was rejected by the separate verifier with expected exit code 2. Failed rule: L015 / F051-F054 / L020. Failed verifier gates: VB-D-EAUR, VERDICTREADY-GATE and EXPORT-GATE. The export boundary also recorded a legal-disposition mismatch at VB-B-MOCC.

VerdictReady NO; export_authorized false; export files opened 0; bytes written 0; final output inventory []. This was a recorded substantive refusal, not an unexplained crash or missing result.

Controls and execution discipline

All seven prescribed altered-copy controls were rejected: reference law, restored cost fact, wrong disposition, Reed in trace, missing export, forged receipt and changed ZIP bytes. The intentionally missing artifact tests completeness only. Oracle independence, baseline recalculation by the new evaluator and forged-disposition export refusal also passed.

The pinned ZIP, its manifest and the evaluator/verifier hashes passed preflight integrity checks. Benchmark steps ran sequentially with a lock and 120-second process-group timeouts; all completed with their expected exits. Commands, stdout/stderr, fresh receipts and environment records are included. Existing evidence stayed unchanged and replay automation stayed paused.

Source and interpretation

Pin: ac3a8cd522718235d2a7154c728566614086a809. The original reference evidence remains reference only. The evaluator was newly implemented for the earlier Demo 3; this activation executed that unchanged implementation. Its verifier is separate internal code. This is not original-engine execution, universal model determinism, unaffiliated validation, external certification or legal compliance.

Byte receipts and scope

Each digest below was measured independently for fresh Runs A, B and C; all three match the pinned expected value and historical bytes. HASHES.md lists all fifteen observations explicitly.

01_canonical_input.json | A = B = C = pinned | 14996 bytes

c53a77a0df2032fc2931bfb1793991d274c98ce1a6d2f7aef5534bc3ae0df561

02_normalized_record.csv | A = B = C = pinned | 14563 bytes

99aa74d84fe9e33cedaf9a4820eb846ce5618f0aa74e8951241ba1347ada93cc

03_gate_matrix.json | A = B = C = pinned | 4181 bytes

242401fedf2c6a01b0cc376e8734cfeaffff180f5d669ae9775e43a95c2014775

04_verdictbridge_assessment.txt | A = B = C = pinned | 2348 bytes

90faa90efa8f79aab956427479788e616f87b0178606d4f2563535cccd5a0378

canonical_run.zip | A = B = C = pinned | 36608 bytes

947314014bd8251f8cfb67d080bd7244d648752e9532aee109b85cf5f836324c

What was frozen; what changed

Inputs, controlling rules, expected outcomes, evaluator/verifier source and canonical serialization were unchanged. Each of the three fresh runs used CPython 3.9.6, Darwin 25.6.0, arm64, with the same recorded configuration. Runtime receipts differ from the pinned receipts only in process ID and start/finish timestamps. Those fields are outside the five-artifact identical-byte boundary.

The historical receipt lacks an OS/interpreter inventory, so identical historical execution conditions cannot be established. Artifact equality to that historical release was measured directly. The new complete package necessarily differs because it includes fresh receipts, verification records and this report. No complete-package identity claim is made.

Limits of the result

Three successful executions establish the observed finite replay, not an always-identical guarantee. No generative model was called. The tested changes cover specified fact and integrity controls; they do not cover every possible material change. Cooperative export refusal is not system-wide enforcement or human shutdown of an AI model.

Review and bindings

Start with verification/measured-results.json, execution/commands.json, environment.json and the fresh replay/ directory. Inspect the verifier and invalid-control evidence. The unsigned replay certificate binds the delivered evidence; SHA256SUMS.txt binds the certificate. The complete ZIP hash is provided separately to avoid self-reference.

Next step: an unaffiliated reviewer authenticates the pinned source/package and reviews or independently reproduces the bounded result in a newly recorded environment. Keep that review distinct from this internal verification. Nothing has been published or sent.